

VermaGourav_Assignment7

January 31, 2021

1 Assignment 7.1

1.1 A

1.1.1 Gourav Verma

```
[23]: import os
import json
from pathlib import Path
import gzip
import hashlib
import shutil
import pandas as pd
import pygeohash
import s3fs

endpoint_url='https://storage.budsc.midwest-datascience.com'
current_dir = Path(os.getcwd()).absolute()
results_dir = current_dir.joinpath('results')

if results_dir.exists():
    shutil.rmtree(results_dir)
results_dir.mkdir(parents=True, exist_ok=True)

def read_jsonl_data():
    s3 = s3fs.S3FileSystem(
        anon=True,
        client_kwargs={
            'endpoint_url': endpoint_url
        }
    )
    src_data_path = 'data/processed/openflights/routes.jsonl.gz'
    with s3.open(src_data_path, 'rb') as f_gz:
        with gzip.open(f_gz, 'rb') as f:
            records = [json.loads(line) for line in f.readlines()]
    return records

def flatten_record(record):
```

```

flat_record = dict()

for key, value in record.items():
    if key in ['airline', 'src_airport', 'dst_airport']:
        if isinstance(value, dict):
            for child_key, child_value in value.items():
                flat_key = '{}_{}'.format(key, child_key)
                flat_record[flat_key] = child_value
        else:
            flat_record[key] = value
    return flat_record

def create_flattened_dataset():
    records = read_jsonl_data()
    parquet_path = results_dir.joinpath('routes-flattened.parquet')
    return pd.DataFrame.from_records([flatten_record(record) for record in
    ↪records])

```

```

[24]: df = create_flattened_dataset()
df['key'] = df['src_airport_iata'].astype(str) + df['dst_airport_iata'].
    ↪astype(str) + df['airline_iata'].astype(str)

```

```

[25]: partitions = (
    ('A', 'A'), ('B', 'B'), ('C', 'D'), ('E', 'F'),
    ('G', 'H'), ('I', 'J'), ('K', 'L'), ('M', 'M'),
    ('N', 'N'), ('O', 'P'), ('Q', 'R'), ('S', 'T'),
    ('U', 'U'), ('V', 'V'), ('W', 'X'), ('Y', 'Z')
)

```

```

[26]: #nan values are causing an issue with key assignment so I am removing them from
    ↪the dataset.
df = df[df['src_airport_iata'].isna() == False]

```

```

[27]: #I took these steps to get the appropriate values for the partitions. I'm sure
    ↪there is an easier way but I kept running into
    ↪roadblocks.

#set kv-key equal to the first letter
df['kv_key'] = df['key'].str[0]

#assign a value from the partitions list of tuples
df['kv_key'] = df['kv_key'].apply(lambda x: [str('-'.join(partition)) for
    ↪partition in partitions if (str(x) >= partition[0]) & (str(x) <=
    ↪partition[1])])

# the result of the previous assignment were lists so here I am converting them
    ↪to strings

```

```
df['kv_key'] = [''.join(partition) for partition in df['kv_key']]

#here i'm replacing the partitions that have the same start and end letter with
↳ a single letter
df['kv_key'] = [partition[0] if partition[0] == partition[2] else partition for
↳ partition in df['kv_key']]
```

```
[28]: df.to_parquet(
        path='results/kv',
        partition_cols=['kv_key']
    )
```

1.2 B

```
[29]: import hashlib

def hash_key(key):
    m = hashlib.sha256()
    m.update(str(key).encode('utf-8'))
    return m.hexdigest()
```

```
[30]: df['hashed'] = df['key'].apply(lambda x: hash_key(x))
df['hash_key'] = df['hashed'].str[0]
```

```
[31]: df.head()
```

```
[31]:
```

	airline_id	airline_name	airline_alias	airline_iata	\
0	410	Aerocondor	ANA All Nippon Airways	2B	
1	410	Aerocondor	ANA All Nippon Airways	2B	
2	410	Aerocondor	ANA All Nippon Airways	2B	
3	410	Aerocondor	ANA All Nippon Airways	2B	
4	410	Aerocondor	ANA All Nippon Airways	2B	

	airline_icao	airline_callsign	airline_country	airline_active	\
0	ARD	AEROCONDOR	Portugal	True	
1	ARD	AEROCONDOR	Portugal	True	
2	ARD	AEROCONDOR	Portugal	True	
3	ARD	AEROCONDOR	Portugal	True	
4	ARD	AEROCONDOR	Portugal	True	

	src_airport_id	src_airport_name	...	dst_airport_id	dst	\
0	2965.0	Sochi International Airport	...		N	
1	2966.0	Astrakhan Airport	...		N	
2	2966.0	Astrakhan Airport	...		N	
3	2968.0	Chelyabinsk Balandino Airport	...		N	
4	2968.0	Chelyabinsk Balandino Airport	...		N	

	dst_airport_tz_id	dst_airport_type	dst_airport_source	codeshare	equipment	\
0	Europe/Moscow	airport	OurAirports	False	[CR2]	
1	Europe/Moscow	airport	OurAirports	False	[CR2]	
2	Europe/Moscow	airport	OurAirports	False	[CR2]	
3	Europe/Moscow	airport	OurAirports	False	[CR2]	
4	Asia/Krasnoyarsk	airport	OurAirports	False	[CR2]	

	key	kv_key	hashed	\
0	AERKZN2B	A	652cdec02010381f175efe499e070c8cbaac1522bac59a...	
1	ASFKZN2B	A	9eea5dd88177f8d835b2bb9cb27fb01268122b635b241a...	
2	ASFMV2B	A	161143856af25bd4475f62c80c19f68936a139f653c1d3...	
3	CEKKZN2B	C-D	39aa99e6ae2757341bede9584473906ef1089e30820c90...	
4	CEKOV2B	C-D	143b3389bce68eea3a13ac26a9c76c1fa583ec2bd26ea8...	

	hash_key
0	6
1	9
2	1
3	3
4	1

[5 rows x 42 columns]

```
[32]: df.to_parquet(
        path='results/hash',
        partition_cols=['hash_key']
    )
```

1.3 C

```
[33]: #get hash for datacenters
datacenters = {}

datacenters['west'] = pygeohash.encode(45.5945645, -121.1786823)
datacenters['central'] = pygeohash.encode(41.1544433, -96.0422378)
datacenters['east'] = pygeohash.encode(39.08344, -77.6497145)

print(datacenters)

{'west': 'c21g6s0rs4c7', 'central': '9z7dnebnj8kb', 'east': 'dqby34cjw922'}
```

```
[34]: #cycle through the datacenter dictionary to assign the closest

def closest_datacenter(latitude, longitude):
    geohash = pygeohash.encode(latitude, longitude)
    dist_dict = {}
    closest_datacenter = ''
```

```

last_distance = None
for key, value in datacenters.items():
    dist = pygeohash.geohash_approximate_distance(str(geohash), str(value))
    dist_dict[key] = dist
    if (last_distance == None) or (dist < last_distance):
        closest_datacenter = key
        last_distance = dist

return closest_datacenter

```

```

[35]: df['datacenter'] = df[['src_airport_latitude', 'src_airport_longitude']].
    ↪ apply(lambda x: closest_datacenter(x[0], x[1]), axis=1)

```

```

[36]: df.to_parquet(
    path='results/geo',
    partition_cols=['datacenter']
)

```

1.4 D

```

[37]: def balance_partitions(keys, num_partitions):
    partitions = []

    #get the ideal number of records in each partition
    partition_size = len(keys) / num_partitions

    #get the count of records for each key
    key_grp_cnts = []
    for key in set(keys):
        occurrences = keys.count(key)
        key_grp_cnts.append(tuple([key, occurrences]))

    key_grp_cnts.sort(key=lambda v: v[0].lower())

    total = 0
    partition_list = []
    #loop through the group counts until you exceed partition_size
    for grp in key_grp_cnts:

        #if the total is 0, then this is the first key in the group
        if total == 0:
            min_grp = grp[0]
            last_group = grp[0]

        #if the incremented total exceeds the ideal partition size, then this_
        ↪ key is the max group and reset the total
        if (total + grp[1]) > partition_size:

```

```

        max_grp = last_group
        partition_list.append(tuple([min_grp, max_grp]))
        last_group = grp[0]
        total=0
    else:
        last_group = grp[0]
        total += grp[1]

    #add last partition
    partition_list.append(tuple([min_grp, last_group]))

    return partition_list

```

```

[38]: #start by using a series from the df above as the list of keys
keys = list(df['airline_name'])
num_partitions=10

```

```

[39]: print(balance_partitions(keys, num_partitions))

```

```

[('40-Mile Air', 'Air Foyle'), ('Air Greenland', 'Amaszonas'), ('Amerijet
International', 'China Eastern Airlines'), ('China SSS', 'Eurowings'), ('Excel
Airways', 'Jet Airways'), ('JetBlue Airways', 'Omni Air International'), ('Onur
Air', 'Shaheen Air International'), ('Shanghai Airlines', 'TransAsia Airways'),
('Transavia Holland', 'UTair-Express'), ('Valuair', 'Zoom Airlines')]

```